

[Lec 04]

Back-propagation and computation graphs

Matrix Gradients for simple Neural Net

Derivative with respect to weight matrix

$$\frac{\partial s}{\partial \mathbf{W}} = \boldsymbol{\delta}^T \mathbf{x}^T$$
$$\begin{matrix} [n \times m] & [n \times 1][1 \times m] \end{matrix}$$

Tips


- Carefully Define Variables
- Use Chain Rule and Shape Convention

Pitfall

- Always use pre-trained word vectors
- Fine-tune it only when size of data is large

Computation Graphs and Back propagation

→ Taking Derivatives using the chain rule



- + : distributes the upstream gradient
- max: routes the upstream gradient
- * : switches the upstream gradient

→ For efficiency, compute all gradients at once

Back-Prop in General Computation Graph

- Fprop: visit nodes in topological order
- Bprop: Recursively apply chain rule along computation graph
 - initialize output gradient = 1
 - visit nodes in reverse order

Regularization

→ prevents overfitting when having lots of features


Vectorization

→ looping over word vectors versus concatenating them all into one large matrix and then multiplying soft-max weights


→ Matrices preferred

Nonlinearities


- sigmoid (logistic)



- tanh & hard tanh
 - rescaled and shifted sigmoid (tanh)



- ReLU (Rectified Linear Unit)
 - Best for building feed-forward deep network
 - variants exist (Leaky Relu, Parametric Relu)



Initialization

- normally must initialize weights to small random values
- biases
 - hidden: Initialize to zero
 - Output: Initialize to optimal value if weights were 0

Optimizers

- SGD works fine
 - But hand tune learning rate
- Adaptive Optimizers
 - Adagrad, RMSprop, **Adam**

Learning Rate

→ start around 0.001

must be order of magnitude (powers of 10)

→ Too Big: may diverge vs Too Small: May not be trained by deadline

→ Better to have decreasing learning rate while training