

```
1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6
7
8 from sklearn.model_selection import KFold,cross_val_score, RepeatedStratifiedKFold,StratifiedK
9 from sklearn.impute import SimpleImputer
10 from sklearn.pipeline import Pipeline
11 from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
12 from sklearn.preprocessing import OneHotEncoder,StandardScaler,PowerTransformer
13 from sklearn.compose import ColumnTransformer
14 from sklearn.pipeline import Pipeline
15 from sklearn.linear_model import LogisticRegression
16 from sklearn.svm import SVC
17 from sklearn.impute import SimpleImputer
18 from sklearn.dummy import DummyClassifier
19 from imblearn.over_sampling import SMOTE
20
21 from sklearn.ensemble import AdaBoostClassifier
22 from sklearn.ensemble import GradientBoostingClassifier
23 from sklearn.ensemble import RandomForestClassifier
24 from sklearn.ensemble import ExtraTreesClassifier
25 from sklearn.neighbors import KNeighborsClassifier
26
27 import optuna
28 from xgboost import XGBClassifier
29 from lightgbm import LGBMClassifier
30 from catboost import CatBoostClassifier
31
32 from sklearn.pipeline import make_pipeline
33 from sklearn.pipeline import Pipeline
34 from sklearn.compose import make_column_transformer
35
36 from sklearn.model_selection import KFold, cross_val_predict, train_test_split,GridSearchCV,cr
37 from sklearn.metrics import accuracy_score,classification_report
38
39 #importing plotly and cufflinks in offline mode
40 import cufflinks as cf
41 import plotly.offline
42 cf.go_offline()
43 cf.set_config_file(offline=False, world_readable=True)
44
45
46 import plotly
47 import plotly.express as px
48 import plotly.graph_objs as go
49 import plotly.offline as py
50 from plotly.offline import iplot
51 from plotly.subplots import make_subplots
52 import plotly.figure_factory as ff
53
54 import missingno as msno
```

```

54 import warnings as warning
55
56 import warnings
57 warnings.filterwarnings("ignore")

```



```

1 pd.set_option('max_columns',100)
2 pd.set_option('max_rows',900)
3
4 pd.set_option('max_colwidth',200)
5
6 df = pd.read_csv('heart.csv')
7 df.head()

```

```
1 df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Age                   918 non-null    int64
 1   Sex                   918 non-null    object
 2   ChestPainType         918 non-null    object
 3   RestingBP             918 non-null    int64
 4   Cholesterol           918 non-null    int64
 5   FastingBS             918 non-null    int64
 6   RestingECG            918 non-null    object
 7   MaxHR                 918 non-null    int64
 8   ExerciseAngina        918 non-null    object
 9   Oldpeak               918 non-null    float64
10   ST_Slope              918 non-null    object
11   HeartDisease          918 non-null    int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB

```

```
1 df.duplicated().sum()
```

```
0
```

```

1 def missing(df):
2     missing_number = df.isnull().sum().sort_values(ascending=False)
3     missing_percent = (df.isnull().sum()/df.isnull().count()).sort_values(ascending=False)

```

```

4     missing_values = pd.concat([missing_number, missing_percent], axis=1, keys=['Missing_Number',
5     return missing_values
6
7 missing(df)

```

```

1 numerical= df.drop(['HeartDisease'], axis=1).select_dtypes('number').columns
2
3 categorical = df.select_dtypes('object').columns
4
5 print(f'Numerical Columns: {df[numerical].columns}')
6 print('\n')
7 print(f'Categorical Columns: {df[categorical].columns}')

```

Numerical Columns: Index(['Age', 'RestingBP', 'Cholesterol', 'FastingBS', 'MaxHR', 'Oldpeak',

Categorical Columns: Index(['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slope',



```
1 df[categorical].nunique()
```

```

Sex                2
ChestPainType      4
RestingECG         3
ExerciseAngina     2
ST_Slope           3
dtype: int64

```

```

1 y = df['HeartDisease']
2 print(f'Percentage of patient had a HeartDisease: {round(y.value_counts(normalize=True)[1]*100, 2)}%')

```

Percentage of patient had a HeartDisease: 55.34 % --> (508 patient)

Percentage of patient did not have a HeartDisease: 44.66 % --> (410 patient)

```
1 df['HeartDisease'].iplot(kind='hist')
```

```
1 df[numerical].describe()
```

```
1 df[numerical].iplot(kind='hist');
```

```
1 df[numerical].iplot(kind='histogram',subplots=True,bins=50)
```

```

1 skew_limit = 0.75 # This is our threshold-limit to evaluate skewness. Overall below abs(1) seer
2 skew_vals = df[numerical].drop('FastingBS', axis=1).skew()
3 skew_cols= skew_vals[abs(skew_vals)> skew_limit].sort_values(ascending=False)
4 skew_cols

Oldpeak      1.022872
dtype: float64

1 numerical= df.select_dtypes('number').columns
2
3
4 matrix = np.triu(df[numerical].corr())
5 fig, ax = plt.subplots(figsize=(14,10))
6 sns.heatmap (df[numerical].corr(), annot=True, fmt= '.2f', vmin=-1, vmax=1, center=0, cmap='c

```

```
1 df[categorical].head()
```

▼ Gender and Heart Disease

```
1 print (f'A female person has a probability of {round(df[df["Sex"]=="F"]["HeartDisease"].mean())
2
3 print()
4
5 print (f'A male person has a probability of {round(df[df["Sex"]=="M"]["HeartDisease"].mean())*1
6
7 print()
8
```

A female person has a probability of 25.91 % have a HeartDisease

A male person has a probability of 63.17 % have a HeartDisease

```
1 fig = px.histogram(df, x="Sex", color="HeartDisease",width=400, height=400)
2 fig.show()
```

▼ Chest Pain Type and Heart Disease

```
1 df.groupby('ChestPainType')['HeartDisease'].mean().sort_values(ascending=False)

ChestPainType
ASY      0.790323
TA       0.434783
NAP      0.354680
ATA      0.138728
Name: HeartDisease, dtype: float64

1 fig = px.histogram(df, x="ChestPainType", color="HeartDisease",width=400, height=400)
2 fig.show()
```

▼ RestingECG and Heart Disease

```
1 df.groupby('RestingECG')['HeartDisease'].mean().sort_values(ascending=False)

RestingECG
ST      0.657303
LVH     0.563830
Normal  0.516304
Name: HeartDisease, dtype: float64

1 fig = px.histogram(df, x="RestingECG", color="HeartDisease",width=400, height=400)
2 fig.show()
```


▼ ExerciseAngina and Heart Disease

```
1 df.groupby('ExerciseAngina')['HeartDisease'].mean().sort_values(ascending=False)
```

```
ExerciseAngina  
Y      0.851752  
N      0.351005  
Name: HeartDisease, dtype: float64
```

```
1 fig = px.histogram(df, x="ExerciseAngina", color="HeartDisease", width=400, height=400)  
2 fig.show()
```

▼ ST_Slope and Heart Disease

```

1 df.groupby('ST_Slope')['HeartDisease'].mean().sort_values(ascending=False)

ST_Slope
Flat      0.828261
Down      0.777778
Up         0.197468
Name: HeartDisease, dtype: float64

1 fig = px.histogram(df, x="ST_Slope", color="HeartDisease",width=400, height=400)
2 fig.show()

```

▼ Overall Insights from the Exploratory Data Analysis

```

1 accuracy =[]
2 model_names =[]
3
4
5 X= df.drop('HeartDisease', axis=1)
6 y= df['HeartDisease']
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8
9 ohe= OneHotEncoder()
10 ct= make_column_transformer((ohe,categorical),remainder='passthrough')
11

```

```

12
13 model = DummyClassifier(strategy='constant', constant=1)
14 pipe = make_pipeline(ct, model)
15 pipe.fit(X_train, y_train)
16 y_pred = pipe.predict(X_test)
17 accuracy.append(round(accuracy_score(y_test, y_pred),4))
18 print (f'model : {model} and accuracy score is : {round(accuracy_score(y_test, y_pred),4)}')
19
20 model_names = ['DummyClassifier']
21 dummy_result_df = pd.DataFrame({'Accuracy':accuracy}, index=model_names)
22 dummy_result_df

```

```

1 accuracy =[]
2 model_names =[]
3
4
5 X= df.drop('HeartDisease', axis=1)
6 y= df['HeartDisease']
7 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8
9 ohe= OneHotEncoder()
10 ct= make_column_transformer((ohe,categorical),remainder='passthrough')
11
12
13 lr = LogisticRegression(solver='liblinear')
14 lda= LinearDiscriminantAnalysis()
15 svm = SVC(gamma='scale')
16 knn = KNeighborsClassifier()
17
18 models = [lr,lda,svm,knn]
19
20 for model in models:
21     pipe = make_pipeline(ct, model)
22     pipe.fit(X_train, y_train)
23     y_pred = pipe.predict(X_test)
24     accuracy.append(round(accuracy_score(y_test, y_pred),4))
25     print (f'model : {model} and accuracy score is : {round(accuracy_score(y_test, y_pred),4)}')
26
27 model_names = ['Logistic','LinearDiscriminant','SVM','KNeighbors']
28 result_df1 = pd.DataFrame({'Accuracy':accuracy}, index=model_names)
29 result_df1

```

```

1  accuracy =[]
2  model_names =[]
3
4
5  X= df.drop('HeartDisease', axis=1)
6  y= df['HeartDisease']
7  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8
9  ohe= OneHotEncoder()
10 s= StandardScaler()
11 ct1= make_column_transformer((ohe,categorical),(s,numerical))
12
13
14 lr = LogisticRegression(solver='liblinear')
15 lda= LinearDiscriminantAnalysis()
16 svm = SVC(gamma='scale')
17 knn = KNeighborsClassifier()
18
19 models = [lr,lda,svm,knn]
20
21 for model in models:
22     pipe = make_pipeline(ct1, model)
23     pipe.fit(X_train, y_train)
24     y_pred = pipe.predict(X_test)
25     accuracy.append(round(accuracy_score(y_test, y_pred),4))
26     print (f'model : {model} and accuracy score is : {round(accuracy_score(y_test, y_pred),4)}')
27
28 model_names = ['Logistic_scl','LinearDiscriminant_scl','SVM_scl','KNeighbors_scl']
29 result_df2 = pd.DataFrame({'Accuracy':accuracy}, index=model_names)
30 result_df2

```

```

1  accuracy =[]
2  model_names =[]
3
4

```

```

5  X= df.drop('HeartDisease', axis=1)
6  y= df['HeartDisease']
7  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8
9  ohe= OneHotEncoder()
10 ct= make_column_transformer((ohe,categorical),remainder='passthrough')
11
12 ada = AdaBoostClassifier(random_state=0)
13 gb = GradientBoostingClassifier(random_state=0)
14 rf = RandomForestClassifier(random_state=0)
15 et= ExtraTreesClassifier(random_state=0)
16
17
18
19 models = [ada,gb,rf,et]
20
21 for model in models:
22     pipe = make_pipeline(ct, model)
23     pipe.fit(X_train, y_train)
24     y_pred = pipe.predict(X_test)
25     accuracy.append(round(accuracy_score(y_test, y_pred),4))
26     print (f'model : {model} and accuracy score is : {round(accuracy_score(y_test, y_pred),4)}')
27
28 model_names = ['Ada','Gradient','Random','ExtraTree']
29 result_df3 = pd.DataFrame({'Accuracy':accuracy}, index=model_names)
30 result_df3

```

```

1  accuracy =[]
2  model_names =[]
3
4
5  X= df.drop('HeartDisease', axis=1)
6  y= df['HeartDisease']
7  X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
8
9  ohe= OneHotEncoder()
10 ct= make_column_transformer((ohe,categorical),remainder='passthrough')
11
12 xgbc = XGBClassifier(random_state=0)
13 lgbmc=LGBMClassifier(random_state=0)
14

```

```

15
16 models = [xgbc, lgbmc]
17
18 for model in models:
19     pipe = make_pipeline(ct, model)
20     pipe.fit(X_train, y_train)
21     y_pred = pipe.predict(X_test)
22     accuracy.append(round(accuracy_score(y_test, y_pred), 4))
23
24 model_names = ['XGBoost', 'LightGBM']
25 result_df4 = pd.DataFrame({'Accuracy': accuracy}, index=model_names)
26 result_df4

```

```

1 accuracy = []
2 model_names = []
3
4
5 X= df.drop('HeartDisease', axis=1)
6 y= df['HeartDisease']
7 categorical_features_indices = np.where(X.dtypes != np.float)[0]
8
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
10
11 model = CatBoostClassifier(verbose=False, random_state=0)
12
13 model.fit(X_train, y_train, cat_features=categorical_features_indices, eval_set=(X_test, y_test))
14 y_pred = model.predict(X_test)
15 accuracy.append(round(accuracy_score(y_test, y_pred), 4))
16
17 model_names = ['Catboost_default']
18 result_df5 = pd.DataFrame({'Accuracy': accuracy}, index=model_names)
19 result_df5
20
21

```

```

1 def objective(trial):
2     X= df.drop('HeartDisease', axis=1)
3     y= df['HeartDisease']
4     categorical_features_indices = np.where(X.dtypes != np.float)[0]
5
6     X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
7
8     param = {

```

```

9      "objective": trial.suggest_categorical("objective", ["Logloss", "CrossEntropy"]),
10     "colsample_bylevel": trial.suggest_float("colsample_bylevel", 0.01, 0.1),
11     "depth": trial.suggest_int("depth", 1, 12),
12     "boosting_type": trial.suggest_categorical("boosting_type", ["Ordered", "Plain"]),
13     "bootstrap_type": trial.suggest_categorical(
14         "bootstrap_type", ["Bayesian", "Bernoulli", "MVS"]
15     ),
16     "used_ram_limit": "3gb",
17 }
18
19 if param["bootstrap_type"] == "Bayesian":
20     param["bagging_temperature"] = trial.suggest_float("bagging_temperature", 0, 10)
21 elif param["bootstrap_type"] == "Bernoulli":
22     param["subsample"] = trial.suggest_float("subsample", 0.1, 1)
23
24 cat_cls = CatBoostClassifier(**param)
25
26 cat_cls.fit(X_train, y_train, eval_set=[(X_test, y_test)], cat_features=categorical_features)
27
28 preds = cat_cls.predict(X_test)
29 pred_labels = np.rint(preds)
30 accuracy = accuracy_score(y_test, pred_labels)
31 return accuracy
32
33
34 if __name__ == "__main__":
35     study = optuna.create_study(direction="maximize")
36     study.optimize(objective, n_trials=50, timeout=600)
37
38     print("Number of finished trials: {}".format(len(study.trials)))
39
40     print("Best trial:")
41     trial = study.best_trial
42
43     print("  Value: {}".format(trial.value))
44
45     print("  Params: ")
46     for key, value in trial.params.items():
47         print("    {}: {}".format(key, value))

```

```

[I 2022-04-04 11:09:46,497] Trial 3 finished with value: 0.8659420289855072 and parameters:
[I 2022-04-04 11:09:49,333] Trial 4 finished with value: 0.8804347826086957 and parameters:
[I 2022-04-04 11:09:55,124] Trial 5 finished with value: 0.8913043478260869 and parameters:
[I 2022-04-04 11:09:56,384] Trial 6 finished with value: 0.8876811594202898 and parameters:
[I 2022-04-04 11:09:59,400] Trial 7 finished with value: 0.8876811594202898 and parameters:
[I 2022-04-04 11:10:01,216] Trial 8 finished with value: 0.8876811594202898 and parameters:
[I 2022-04-04 11:10:02,627] Trial 9 finished with value: 0.8913043478260869 and parameters:
[I 2022-04-04 11:10:06,275] Trial 10 finished with value: 0.8913043478260869 and parameters:
[I 2022-04-04 11:10:07,608] Trial 11 finished with value: 0.894927536231884 and parameters:
[I 2022-04-04 11:10:08,988] Trial 12 finished with value: 0.8840579710144928 and parameters:
[I 2022-04-04 11:10:10,432] Trial 13 finished with value: 0.8876811594202898 and parameters:
[I 2022-04-04 11:10:11,768] Trial 14 finished with value: 0.8876811594202898 and parameters:

[I 2022-04-04 11:10:13,133] Trial 15 finished with value: 0.8913043478260869 and parameters:
[I 2022-04-04 11:10:14,787] Trial 16 finished with value: 0.8876811594202898 and parameters:
[I 2022-04-04 11:10:16,118] Trial 17 finished with value: 0.8840579710144928 and parameters:
[I 2022-04-04 11:10:17,620] Trial 18 finished with value: 0.8913043478260869 and parameters:
[I 2022-04-04 11:10:20,256] Trial 19 finished with value: 0.8768115942028986 and parameters:

```

```

[ I 2022-04-04 11:10:21,998] Trial 20 finished with value: 0.8804347826086957 and parameter:
[ I 2022-04-04 11:10:23,460] Trial 21 finished with value: 0.8840579710144928 and parameter:
[ I 2022-04-04 11:10:24,807] Trial 22 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:26,146] Trial 23 finished with value: 0.8876811594202898 and parameter:
[ I 2022-04-04 11:10:27,444] Trial 24 finished with value: 0.9021739130434783 and parameter:
[ I 2022-04-04 11:10:28,674] Trial 25 finished with value: 0.8876811594202898 and parameter:
[ I 2022-04-04 11:10:29,414] Trial 26 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:34,060] Trial 27 finished with value: 0.8913043478260869 and parameter:
[ I 2022-04-04 11:10:35,374] Trial 28 finished with value: 0.8876811594202898 and parameter:
[ I 2022-04-04 11:10:36,093] Trial 29 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:36,786] Trial 30 finished with value: 0.8840579710144928 and parameter:
[ I 2022-04-04 11:10:37,530] Trial 31 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:38,303] Trial 32 finished with value: 0.8985507246376812 and parameter:
[ I 2022-04-04 11:10:39,083] Trial 33 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:39,760] Trial 34 finished with value: 0.8768115942028986 and parameter:
[ I 2022-04-04 11:10:41,120] Trial 35 finished with value: 0.8804347826086957 and parameter:
[ I 2022-04-04 11:10:41,925] Trial 36 finished with value: 0.8913043478260869 and parameter:
[ I 2022-04-04 11:10:43,250] Trial 37 finished with value: 0.8913043478260869 and parameter:
[ I 2022-04-04 11:10:44,048] Trial 38 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:44,767] Trial 39 finished with value: 0.8913043478260869 and parameter:
[ I 2022-04-04 11:10:45,545] Trial 40 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:46,383] Trial 41 finished with value: 0.8913043478260869 and parameter:
[ I 2022-04-04 11:10:47,223] Trial 42 finished with value: 0.8985507246376812 and parameter:
[ I 2022-04-04 11:10:48,170] Trial 43 finished with value: 0.8876811594202898 and parameter:
[ I 2022-04-04 11:10:48,920] Trial 44 finished with value: 0.8804347826086957 and parameter:
[ I 2022-04-04 11:10:49,806] Trial 45 finished with value: 0.8913043478260869 and parameter:
[ I 2022-04-04 11:10:50,563] Trial 46 finished with value: 0.8985507246376812 and parameter:
[ I 2022-04-04 11:10:51,341] Trial 47 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:52,220] Trial 48 finished with value: 0.894927536231884 and parameter:
[ I 2022-04-04 11:10:53,107] Trial 49 finished with value: 0.8876811594202898 and parameter:

```

Number of finished trials: 50

Best trial:

Value: 0.9021739130434783

Params:

objective: Logloss

colsample_bylevel: 0.01134943888867012

depth: 9

boosting_type: Ordered

bootstrap_type: MVS

```

1 accuracy =[]
2 model_names =[]
3
4
5 X= df.drop('HeartDisease', axis=1)
6 y= df['HeartDisease']
7 categorical_features_indices = np.where(X.dtypes != np.float)[0]
8
9 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
10
11 model = CatBoostClassifier(verbose=False,random_state=0,
12                             objective= 'CrossEntropy',
13                             colsample_bylevel= 0.04292240490294766,
14                             depth= 10,
15                             boosting_type= 'Plain',
16                             bootstrap_type= 'MVS')

```



```
17
18 model.fit(X_train, y_train, cat_features=categorical_features_indices, eval_set=(X_test, y_test))
19 y_pred = model.predict(X_test)
20 accuracy.append(round(accuracy_score(y_test, y_pred), 4))
21 print(classification_report(y_test, y_pred))
22
23 model_names = ['Catboost_tuned']
24 result_df6 = pd.DataFrame({'Accuracy': accuracy}, index=model_names)
25 result_df6
26
27
```

```
1 feature_importance = np.array(model.get_feature_importance())
2 features = np.array(X_train.columns)
3 fi={'features': features, 'feature_importance': feature_importance}
4 df-fi = pd.DataFrame(fi)
5 df-fi.sort_values(by=['feature_importance'], ascending=True, inplace=True)
6 fig = px.bar(df-fi, x='feature_importance', y='features', title="CatBoost Feature Importance", height=400)
7 fig.show()
```

```
1 result_final = pd.concat([dummy_result_df,result_df1,result_df2,result_df3,result_df4,result_d

1  result_final.sort_values(by=['Accuracy'], ascending=True,inplace=True)
2  fig = px.bar(result_final, x='Accuracy', y=result_final.index,title='Model Comparison',height=
3  fig.show()
```

✓ 0초 오후 8:10에 완료됨

